



KubeCon



CloudNativeCon

North America 2019

Specialized Network Protocols for IoT+Edge with Kubernetes

Steven Wong, VMware

Dejan Bosanac, Red Hat



Abstract

This session will survey communication protocols and technologies used in the edge and IoT space.

These use cases can call for specialized protocols and transports:

publish subscribe, multicast

protocols tolerant of intermittent connectivity

Protocols popular in industry verticals (vehicle bus, industrial automation, building automation)

In some cases, support exists now for use with Kubernetes. If not, device gateways and protocol converters might be an option.

Agenda (Intro):

- survey of protocols and transport standards for IoT and edge
- Intro to how a device gateway or protocol converter works
- Intro to extending Kubernetes with CRDs to manage new device types

Agenda (Deep Dive):

- Futures: Could the service mesh concept be extended beyond TCP, HTTP(s)?
- Demonstration: Kubernetes management of an edge application using a specialized protocol
- Demonstration: Use a device gateway with Kubernetes

Speakers



Dejan Bosanac
Red Hat
@dejanb



Steve Wong
VMware
@cantbewong



KubeCon



CloudNativeCon

North America 2019

Agenda

Part 1: Intro

- survey of protocols and transport standards for IoT and edge
- device gateways & protocol converter
- Intro to extending Kubernetes with CRDs to manage new device types

Intermission (5 minutes): Meet others, “birds of a feather”

Part 2: Deep Dive

- Futures: Could the service mesh concept be extended beyond TCP, HTTP(s)?
- Demonstration: Kubernetes management of an edge application using a specialized protocol
- Demonstration: Use a device gateway with Kubernetes

How to get involved with the IoT Edge Working Group



KubeCon



CloudNativeCon

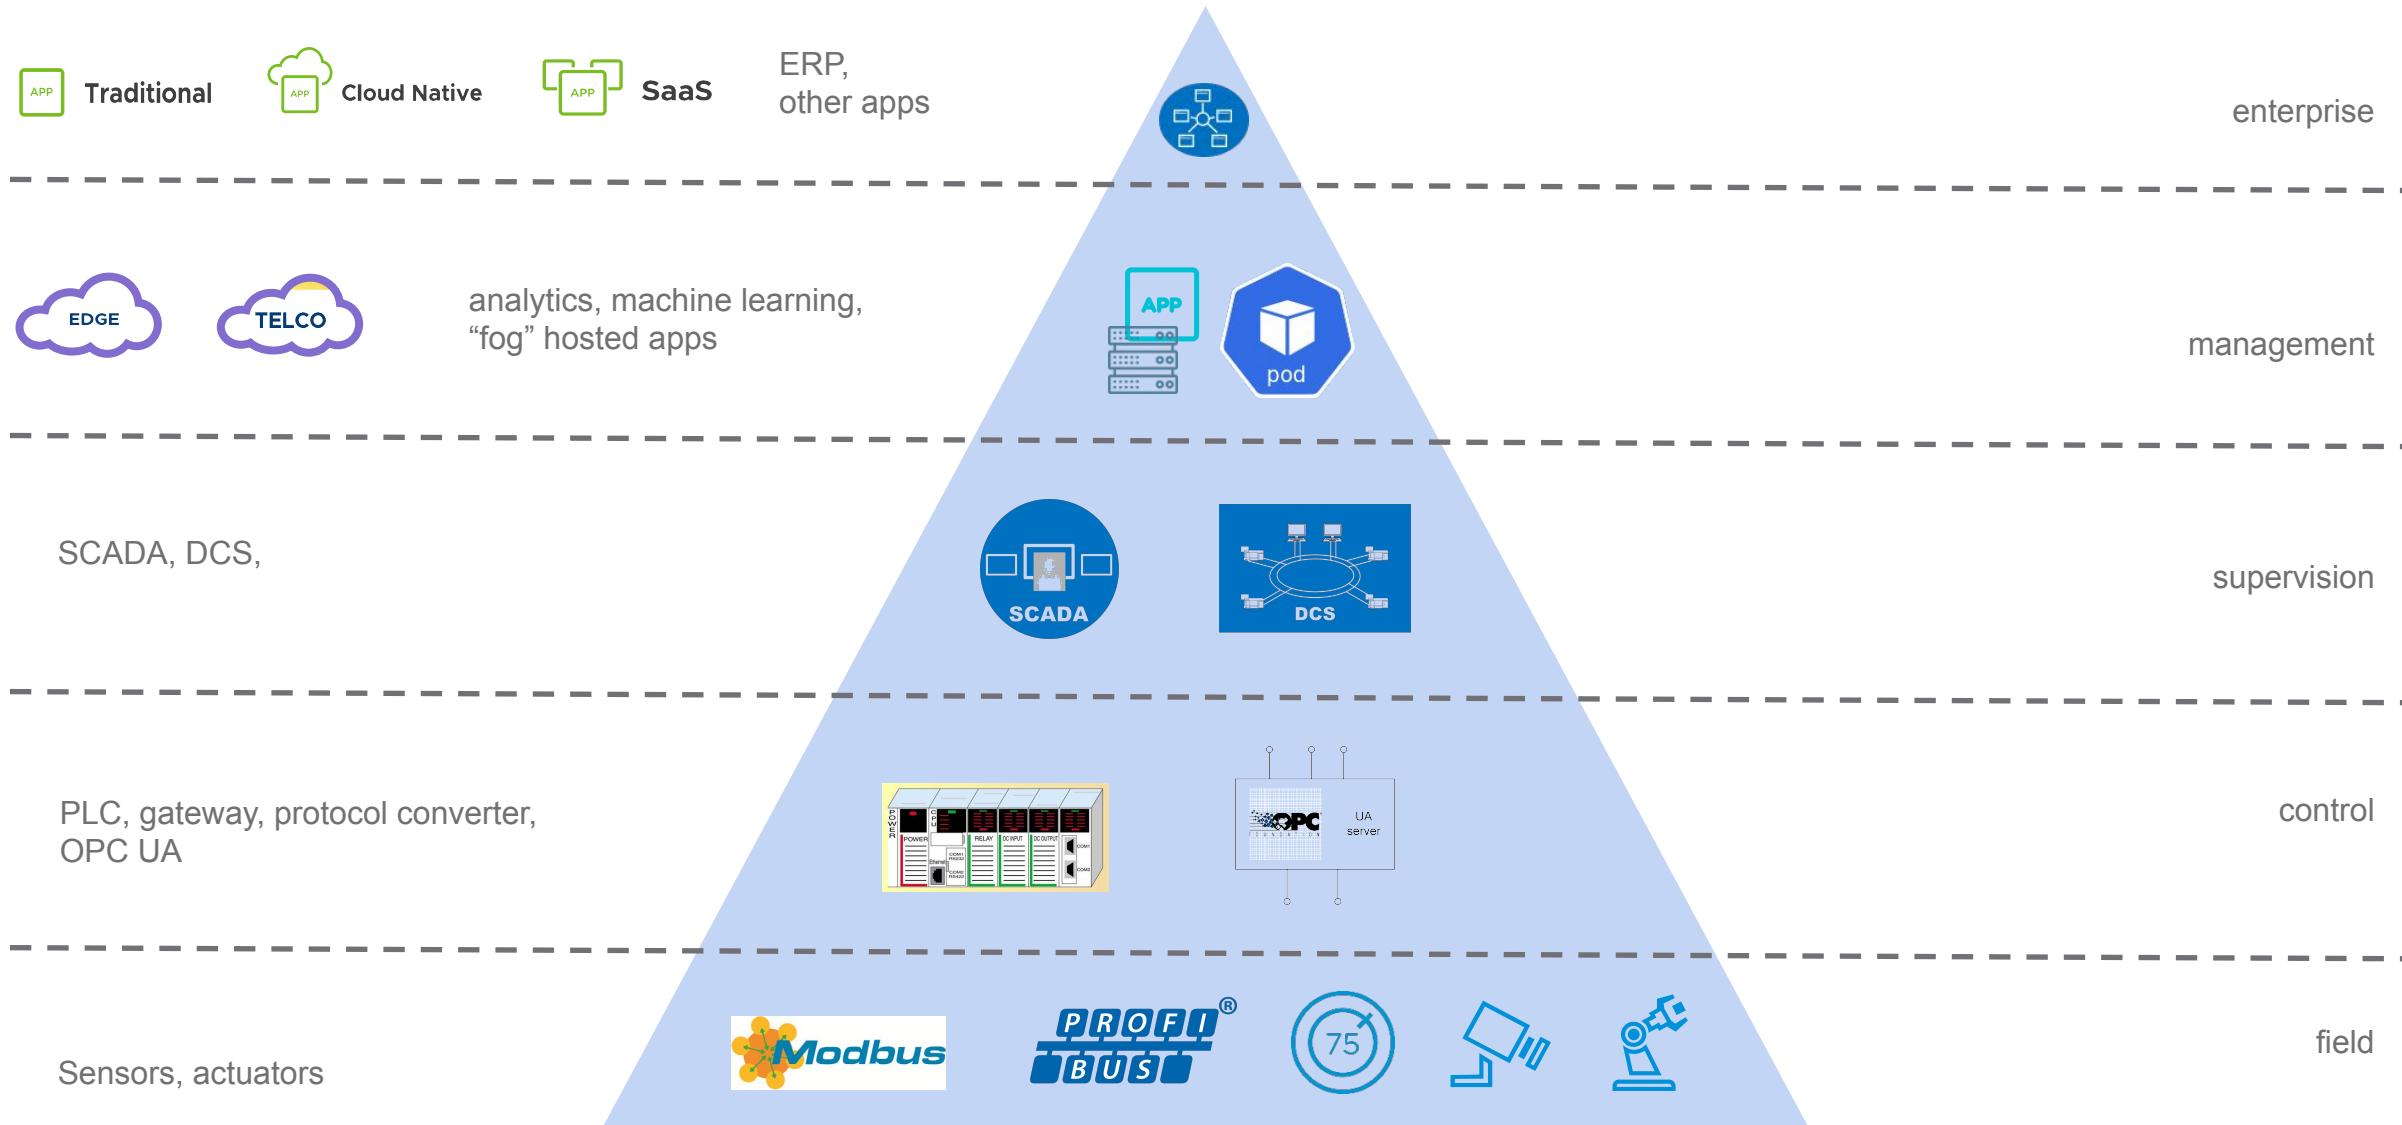
North America 2019

Intro

Survey of protocols and transport standards for
IoT and edge

Applications are multi-level - maybe at multiple locations

Tiers of interconnects support this, often using specialized protocols



Data communication operates at multiple levels

Not unusual to have all these going on simultaneously

May be multiple layers of protocols
(OSI model)



Hardware / physical media: examples:
ethernet, EIA-485 multidrop, EIA-232, etc.

Data Link, Network, Transport – example
MAC, IP, TCP, UDP

Application – example http, domain specific
APIs on http

data and control may be separated,
data may be republished in a transformed
context

Data Plane – low level “raw” data flows

Control Plane

- **Onboard devices or services**
- **Monitor**
- **Manage**
- **Secure**

Content Plane – data transformed to a different context
(ETL)

Type: built in support for specialized applications

Sync req-response, pub-sub, both, or higher level data exchange agnostic that abstracts low level transport.
Low level transport may be connection based or RPC style

Implementation resource demands – suitability for constrained environments? is a broken/router needed?

Behavior with lossy/unreliable networks – latency limits, QoS support, order & delivery guarantees

Security

Topology: point to point, bus, routable

Support for discovery of nodes and data content

Health of community – standards and certification org, scope and openness of software and hardware platforms, popularity within an application domain

Stability and maturity

Protocols

General Purpose



KubeCon



CloudNativeCon

North America 2019

protocol	standard	info	req-reply	pub-sub	QoS	
HTTP	IETF	Wikipedia	✓		no	
HTTP/2	IETF	Wikipedia	✓	✓	no	
AMQP	OASIS	Wikipedia	✓	✓	yes	amqp.org
MQTT	OASIS	Wikipedia		✓	yes	mqtt.org
COAP	IETF	Wikipedia	✓	✓	yes	coap.technology/
DDS	OMG	Wikipedia		✓	yes	
Kafka		Wikipedia		✓	no	kafka.apache.org

Protocols

Higher Level

protocol	standard	info	transport
LWM2M	OMA	Wikipedia	COAP, on UDP or SMS
OGC SensorThings API	OGC	Wikipedia	CoAP, MQTT, HTTP, 6LowPAN
PPMP	eclipse	Wikipedia	
One2M	one2m.org	Wikipedia	

Specialized Protocols

Industrial, Building Automation

protocol	standard	info
Modbus	<u>Modbus Organization</u>	<u>Wikipedia</u>
BACnet	<u>bacnet.org</u>	<u>Wikipedia</u>
OPC UA	<u>OPC Foundation</u>	<u>Wikipedia</u>

Specialized Protocols

Transportation (Automotive, Avionics, Rail)



KubeCon

CloudNativeCon

North America 2019

protocol	standard	info
A ² B	proprietary	link
AFDX	proprietary	Wikipedia
ARINC 429	ARINC	Wikipedia
Byteflight	byteflight	Wikipedia
CAN	ISO	Wikipedia
D2B	IEC 61030	Wikipedia
IDB-1394	IEEE, 1394 trade assoc	Wikipedia
IEBus	proprietary	Wikipedia
I ² C	proprietary	Wikipedia
ISO 9141-1/-2	ISO	Wikipedia

protocol	standard	info
J1708,J1587	SAE	Wikipedia
J1850	SAE	Wikipedia
J1939, ISO 11783	ISO	Wikipedia
Keyword protocol 2000	ISO	Wikipedia
LIN	ISO	Wikipedia
MOST	proprietary	Wikipedia
Multifunction Vehicle Bus	IEC	Wikipedia
SPI	defacto	Wikipedia
VAN	proprietary	Wikipedia



KubeCon



CloudNativeCon

North America 2019

Device Gateways and Protocols Converters

Why use a Gateway / Protocol convertor



KubeCon

CloudNativeCon

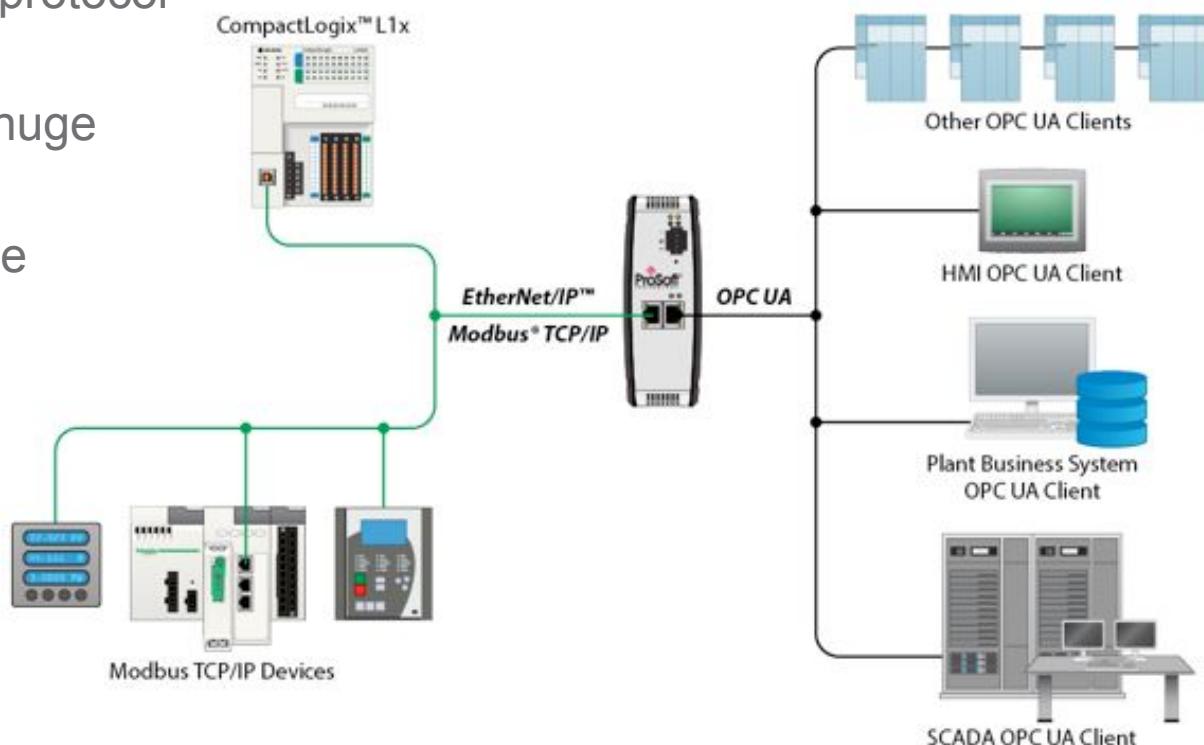
North America 2019

Convert point to point, and local bus protocols to a protocol that is routable

Allow efficient sharing of data sources across multiple consumers

- Low level devices often lack bandwidth and compute to respond to multiple consumers
- Higher tier consumers can be written to utilize just one protocol
 - better than alternative of building and maintaining huge libraries of device drivers in *each* consumer.
 - Better than putting physical media NICs into multiple consumers

Picture source opcfoundation.org



Why use a Gateway / Protocol convertor

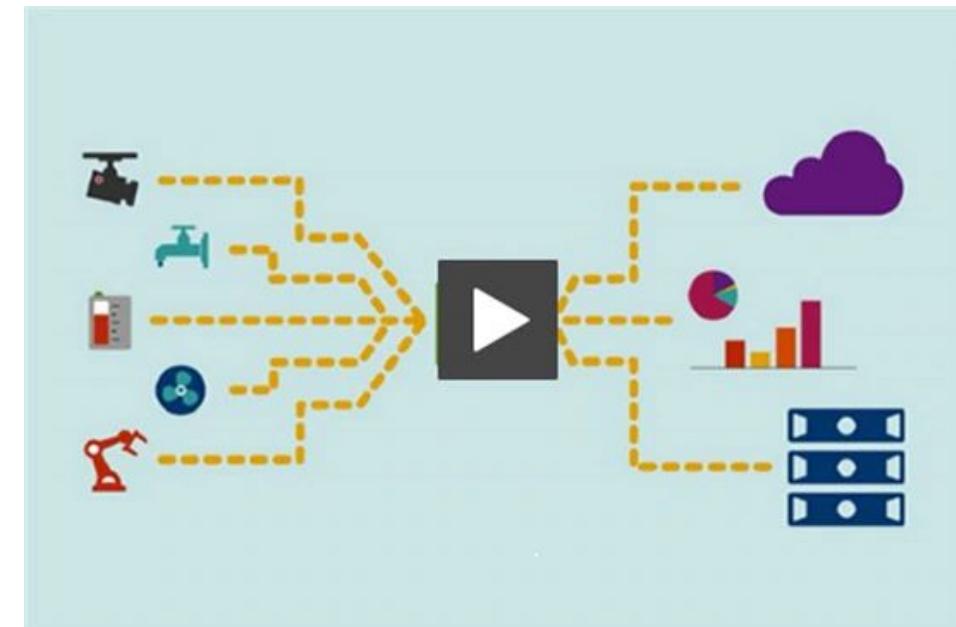
continued

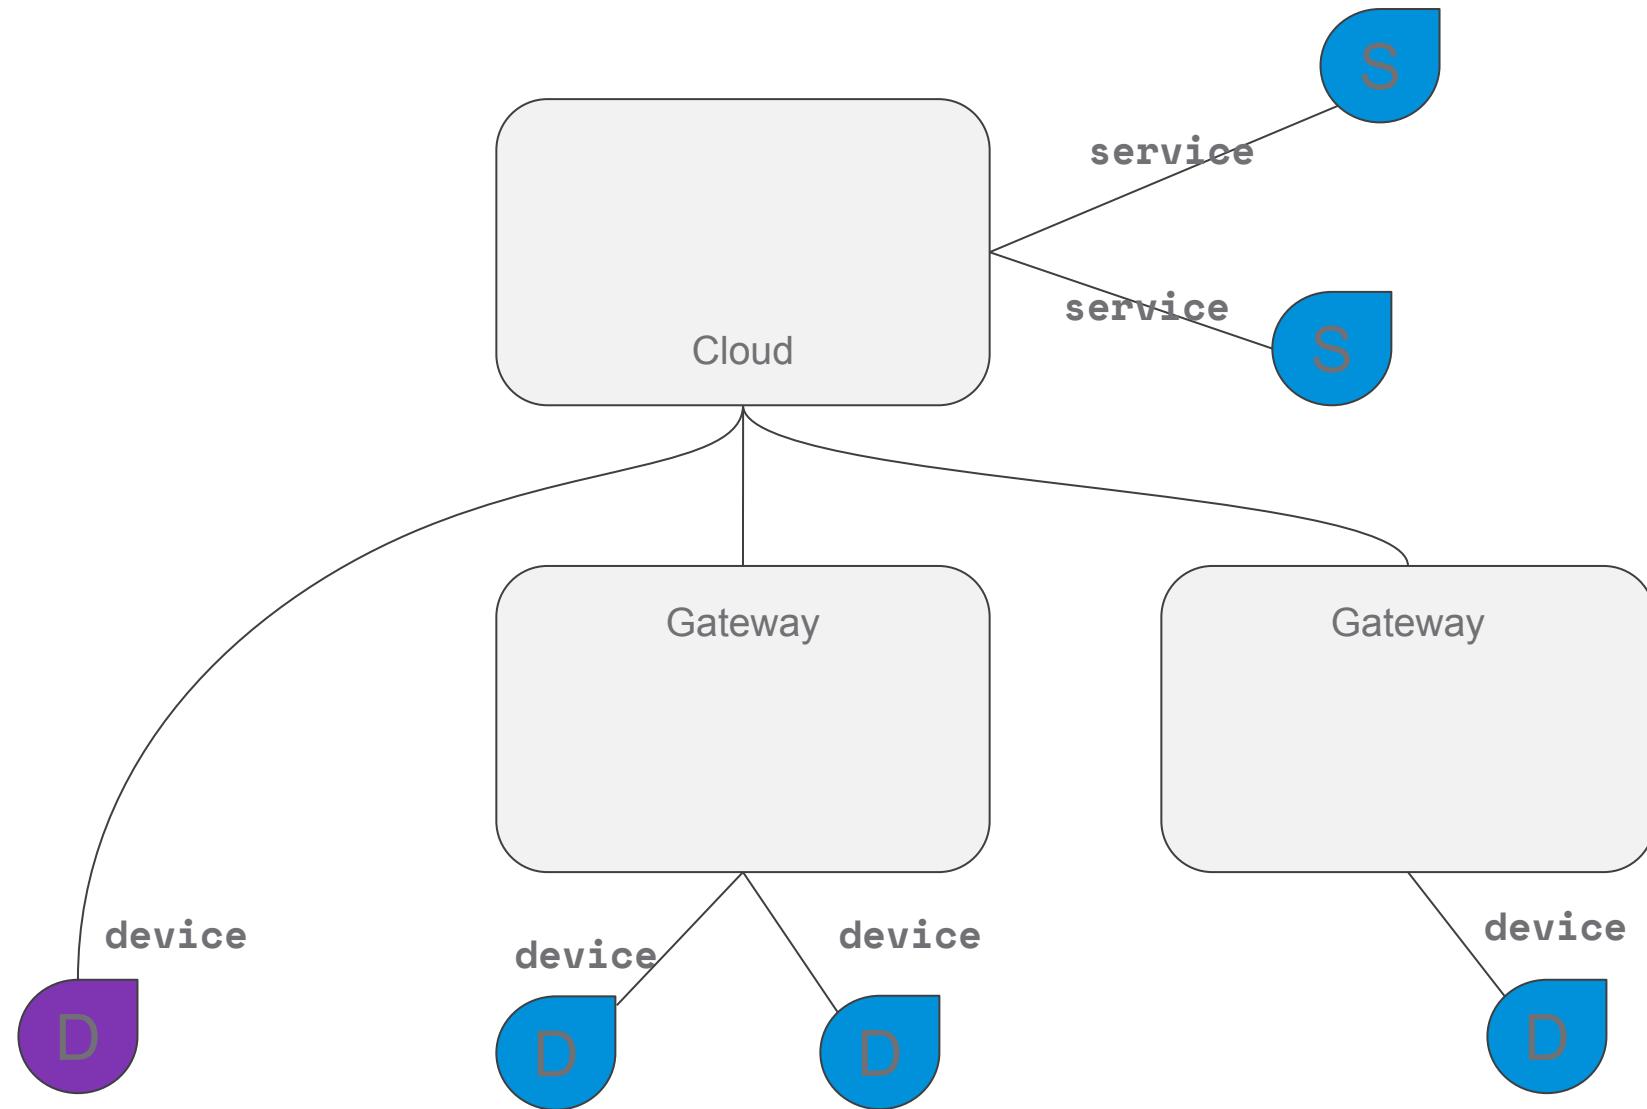
Enable load balancing and monitoring

Can be an injection point to add features to legacy or low level protocols

- Attach timestamps, location, other metadata
- Impose Security

Potentially enable base tier devices to interact with each other, while hiding details and operations from higher tiers.





Rethinking IoT gateways

- Containerization
- Adopt Cloud-native development practices
 - CI/CD
 - Gitops
 - ...

Enabling new use cases

- More resources on the Edge = new use cases
- Machine learning
- Store and forward
- Caching
- ...

Tools for field protocols

- Hardware abstraction layer
- Bluetooth REST API (REST-BLUETE)

IoT cloud connectivity

- Connection oriented
 - MQTT
- RPC style
 - HTTP
- Telemetry
- Command and Control

Eclipse Hono

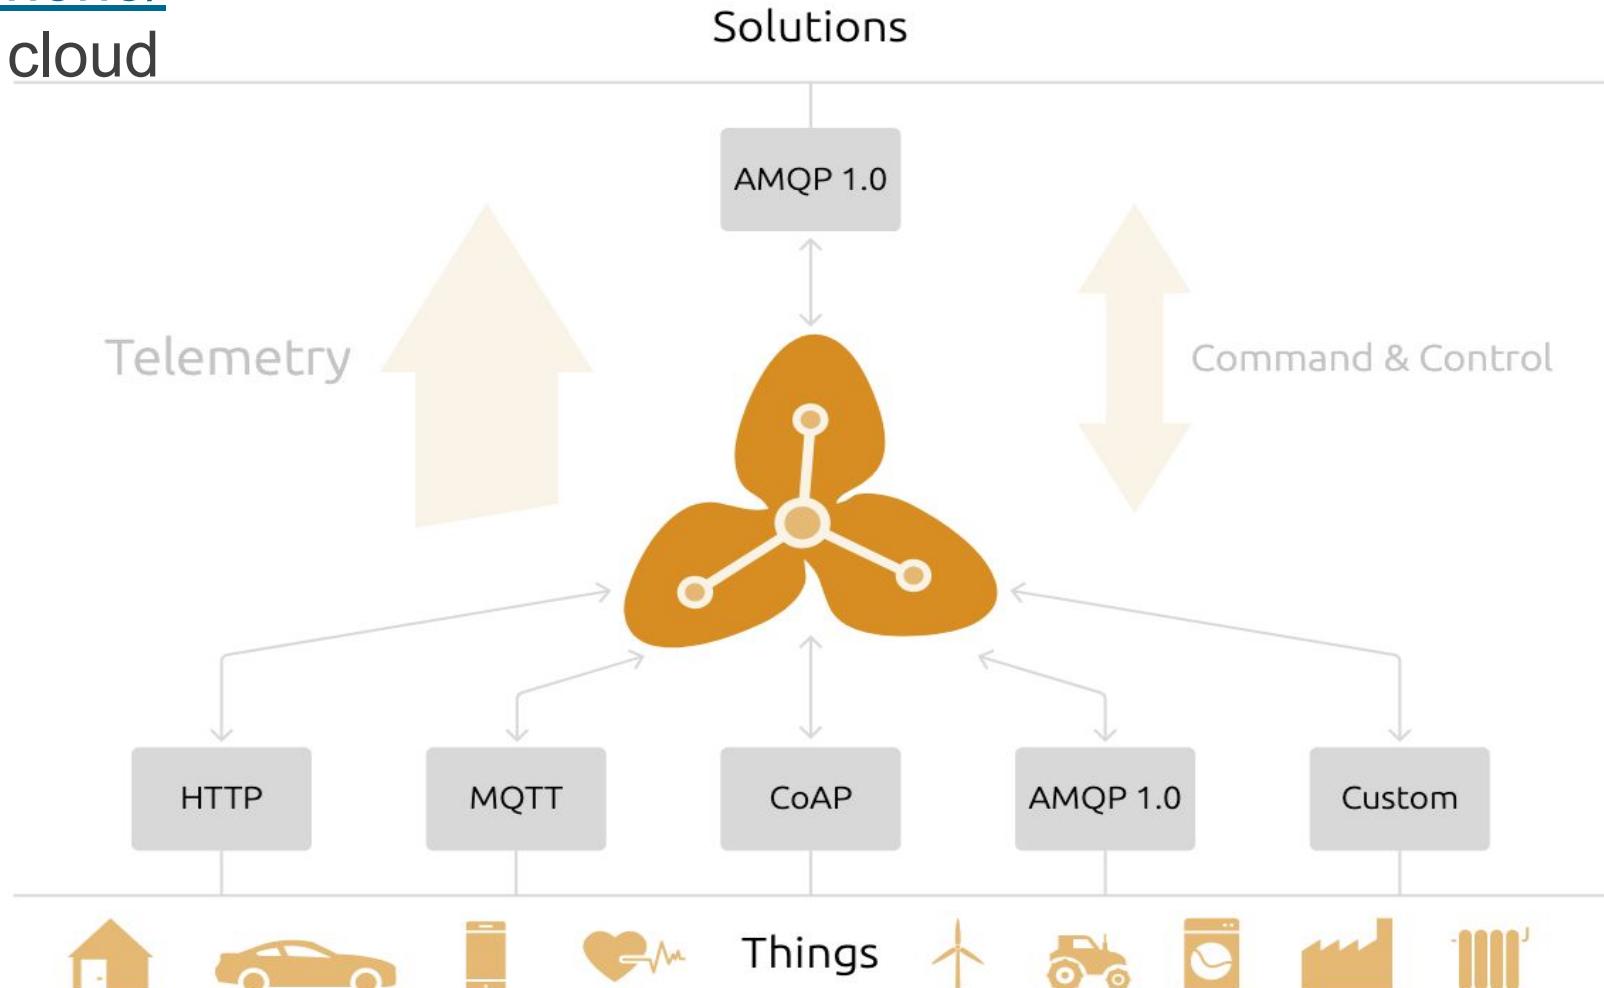


KubeCon

CloudNativeCon

North America 2019

- <https://www.eclipse.org/hono/>
- IoT Connectivity for the cloud
- Scalable
- Secure
- Multi-protocol
- K8s based



Eclipse Hono

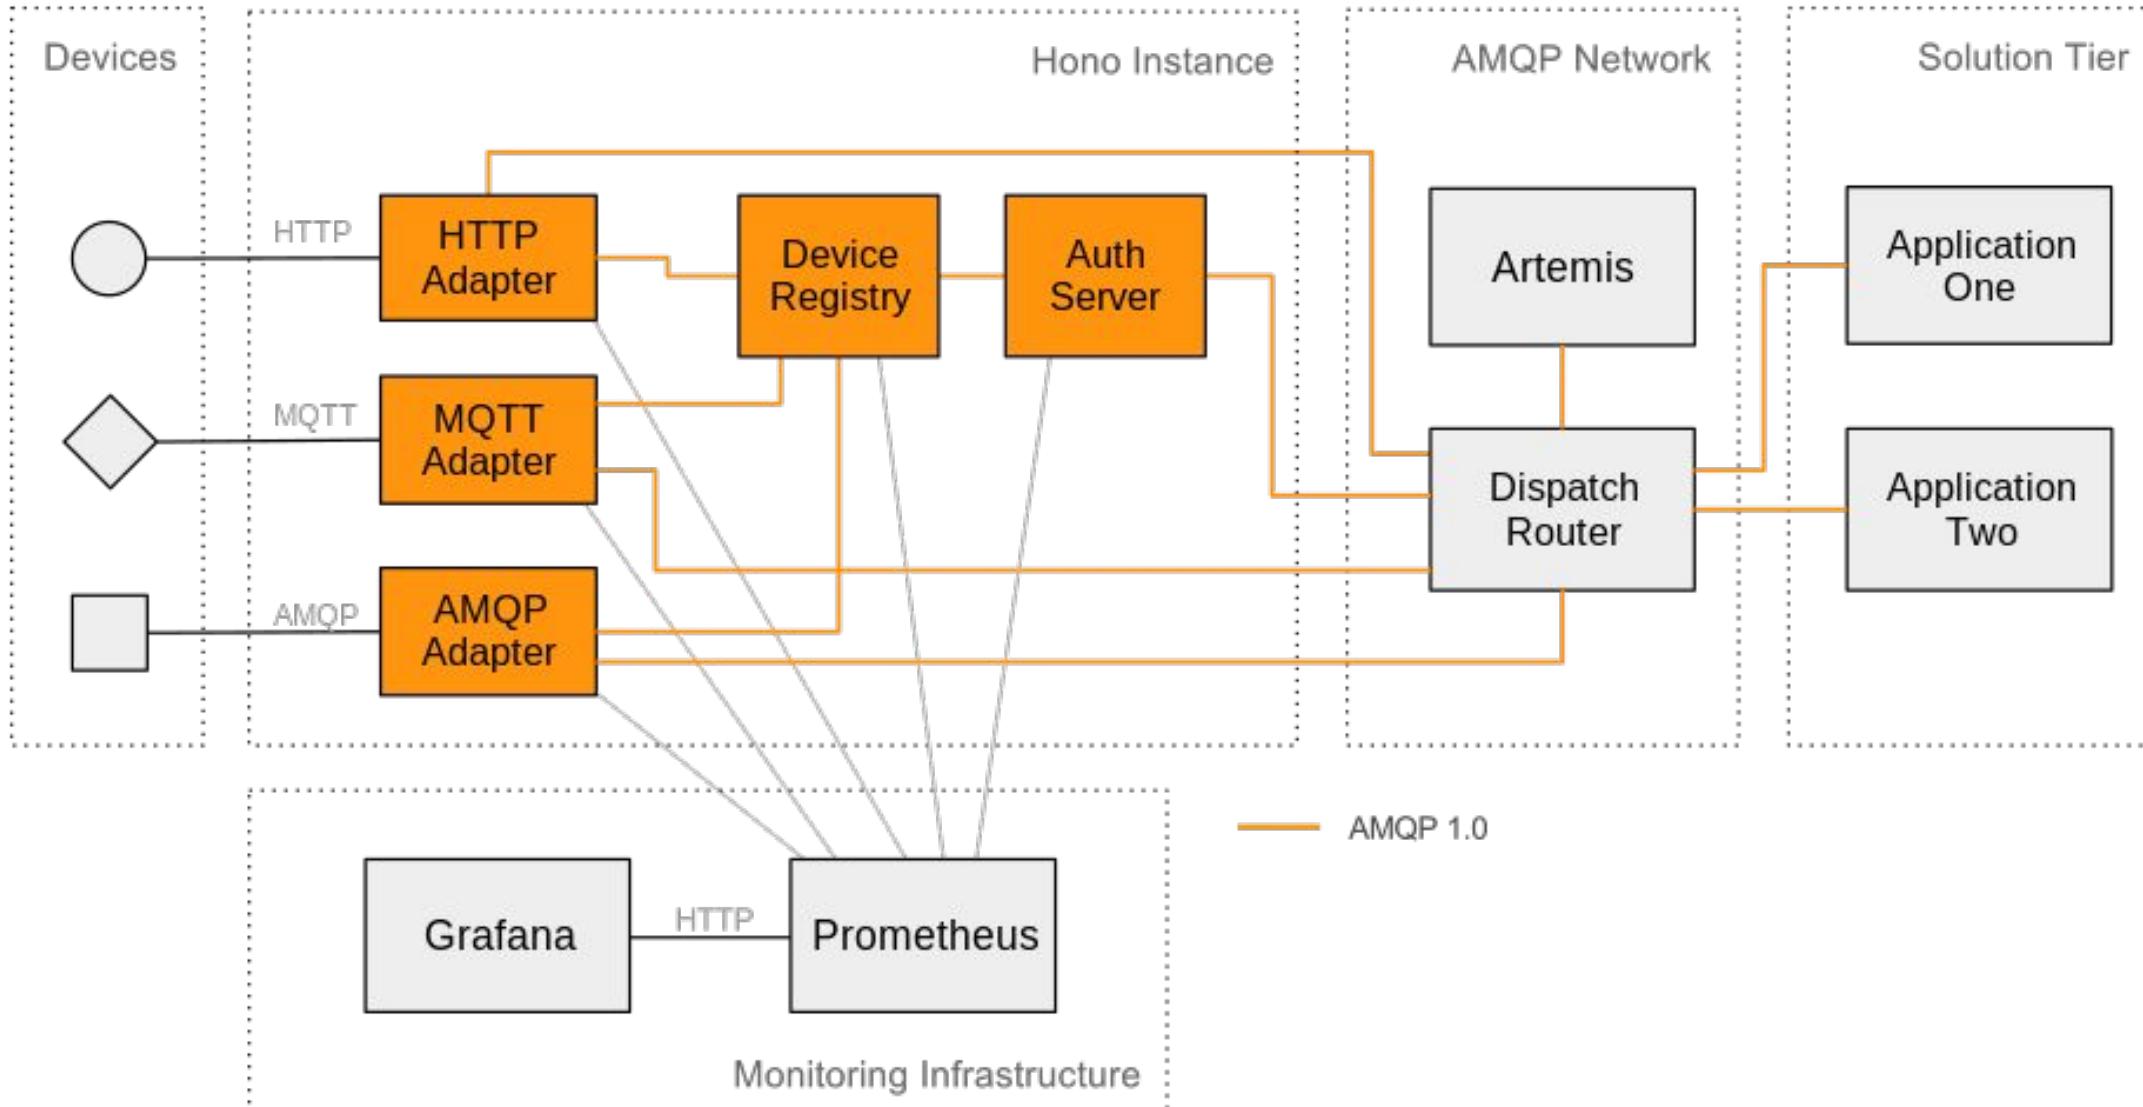


KubeCon



CloudNativeCon

North America 2019



Telemetry

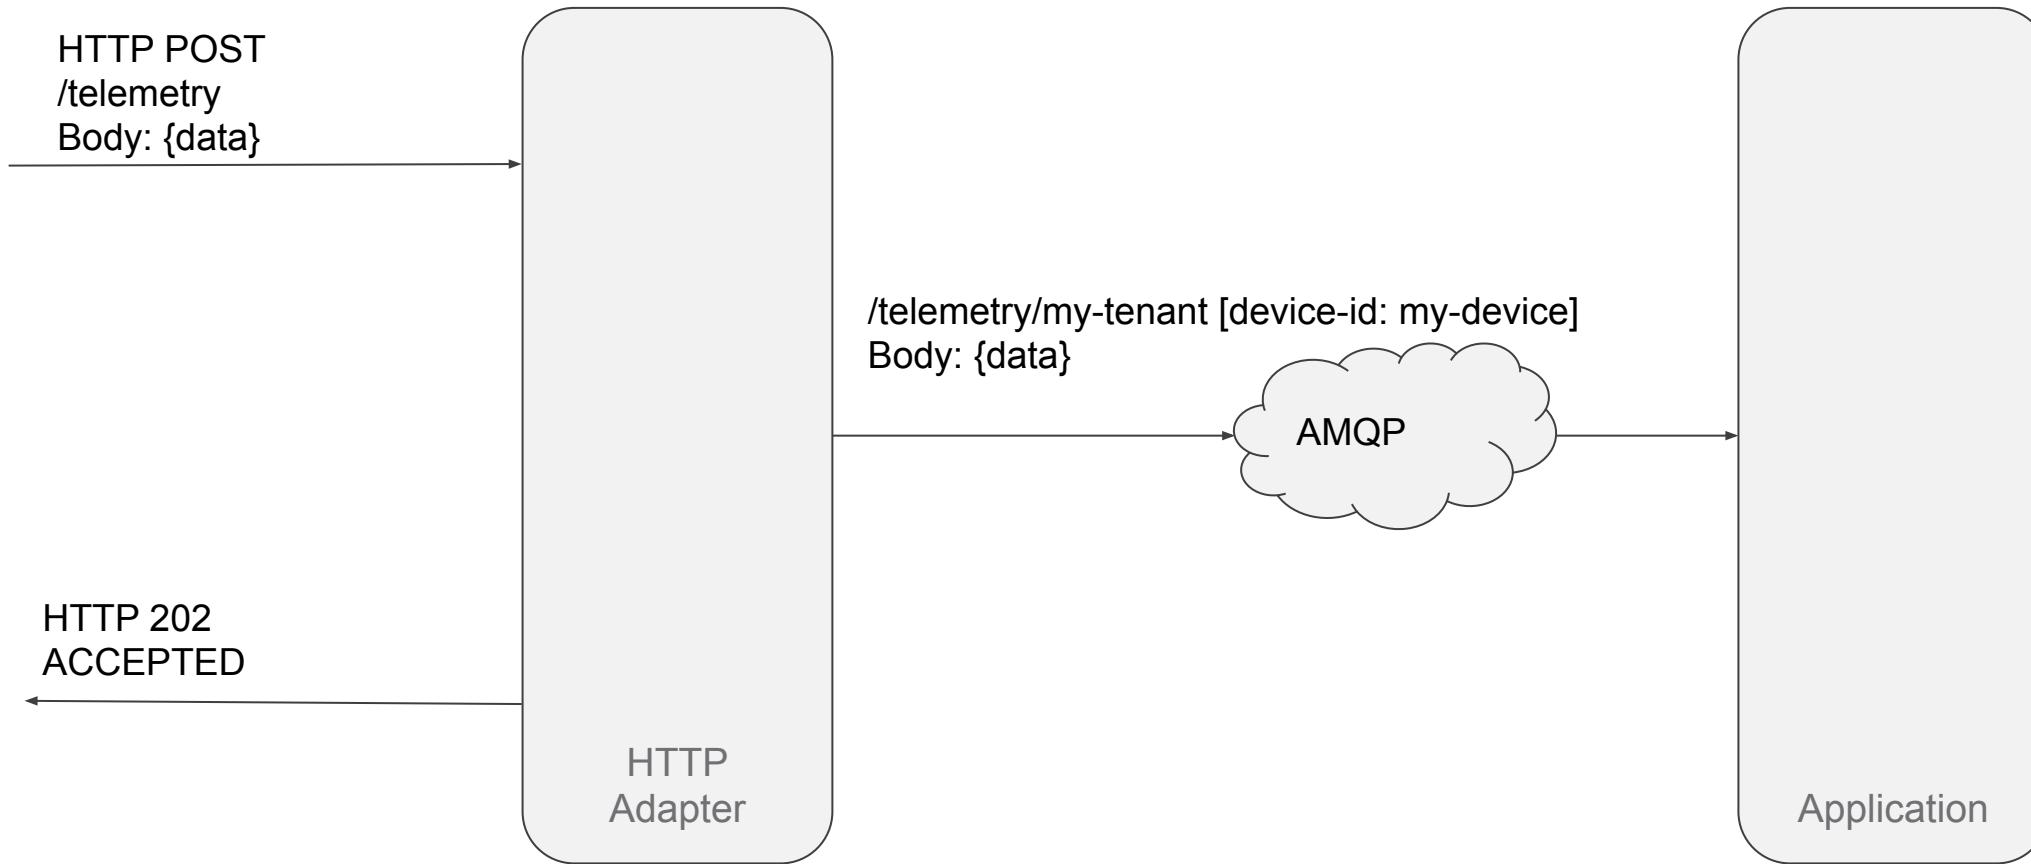


KubeCon



CloudNativeCon

North America 2019



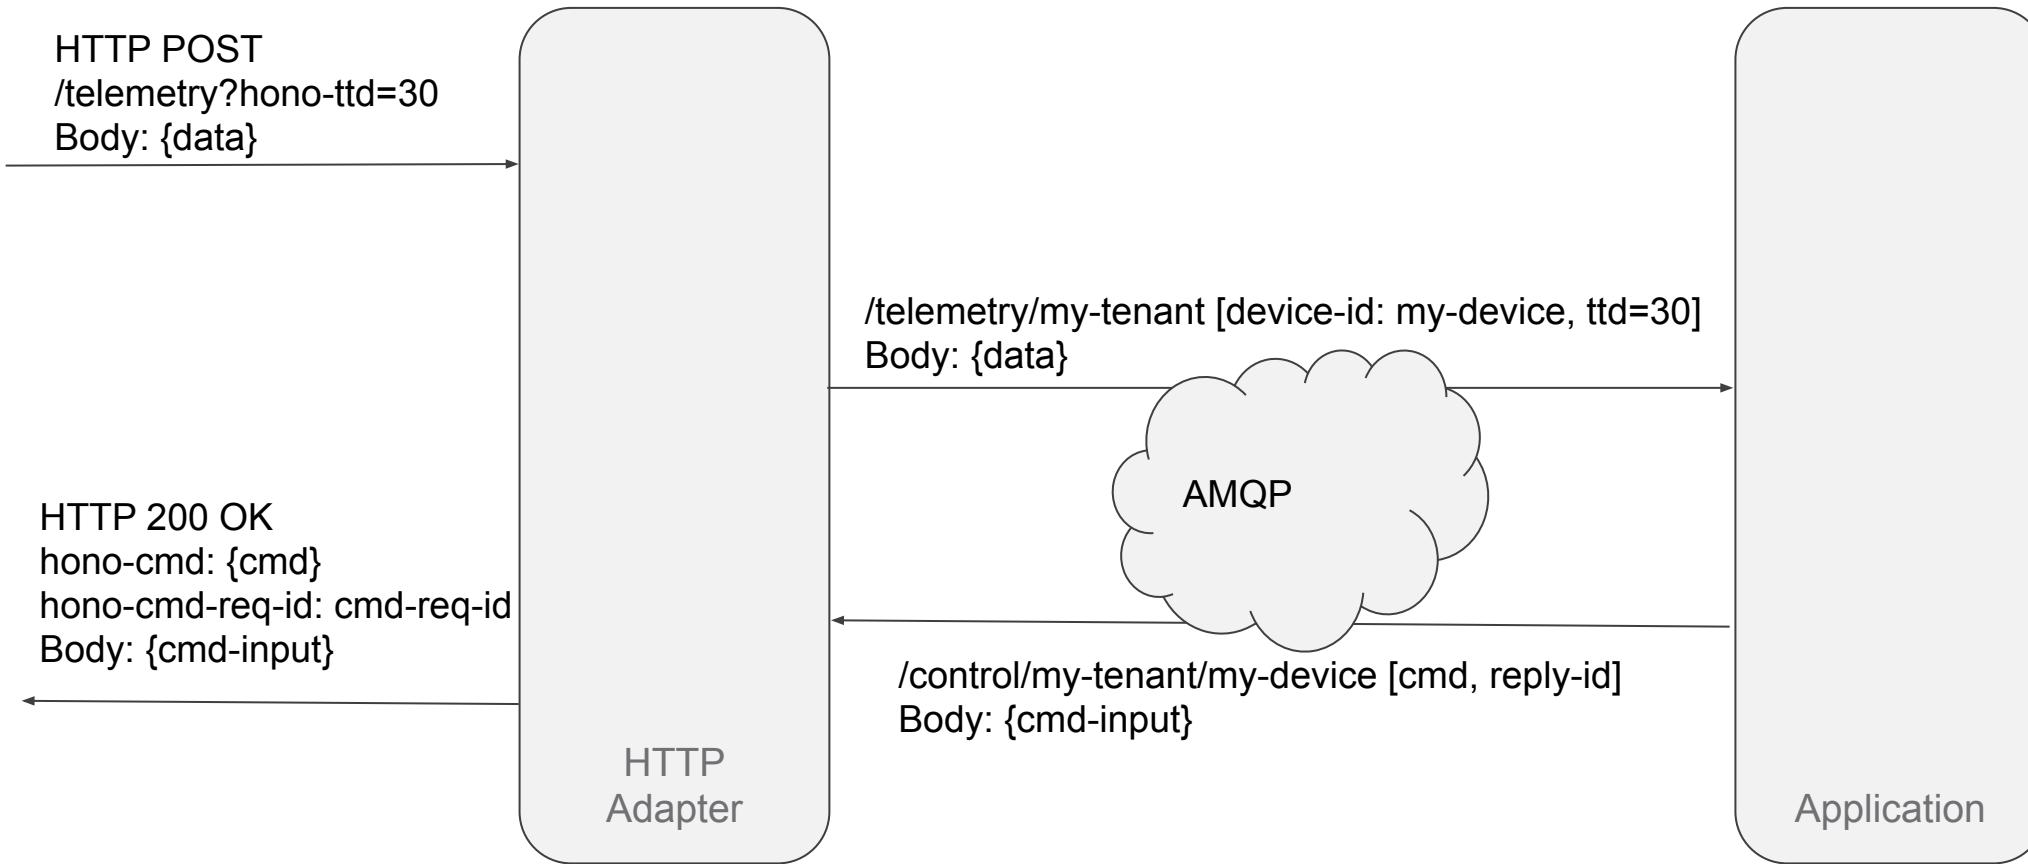
Command and control



KubeCon

CloudNativeCon

North America 2019





KubeCon



CloudNativeCon

North America 2019

CRDs to manage devices?

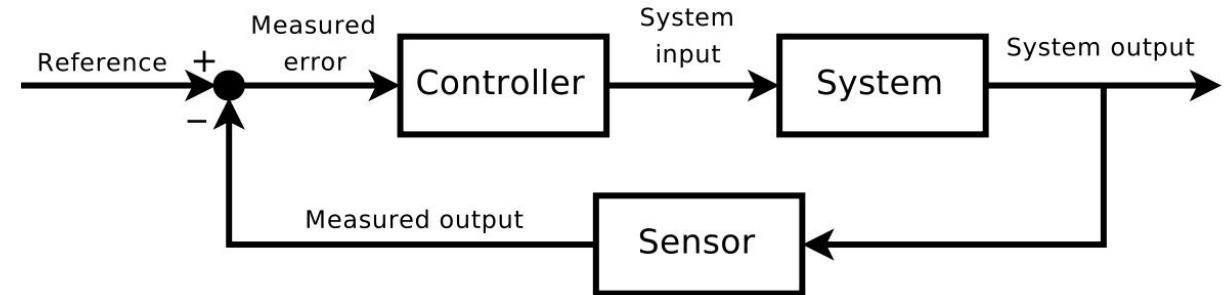
Kubernetes

Based on control loops

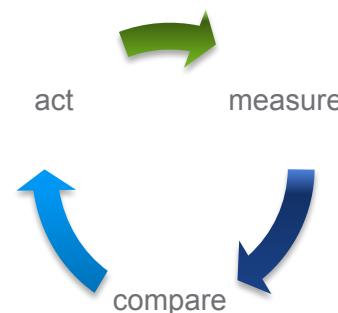
You tell Kubernetes the desired state

Kubernetes relentlessly:

- measures current state
- drives **current state -> desired state**



Recurring pattern of aspects in the system: measure - compare - act



Kubernetes API Server

A REST interface to the etcd database



The API server manages CRUD operations on *resources* like Pods/Deployments/Services

Object properties:

- API version
- Kind
- Metadata
- Specs

The API server itself doesn't actually understand the build-in objects, they might as well be apples, oranges and bananas

Custom Resource Definitions (CRDs) allow you to add *new* resources - allowing management via the Kubernetes API and the kubectl CLI

IoT Sensor

What if we managed via a CRD?

Custom Resource

hamburg-sensor.yaml

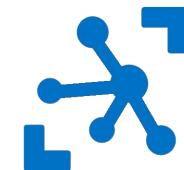
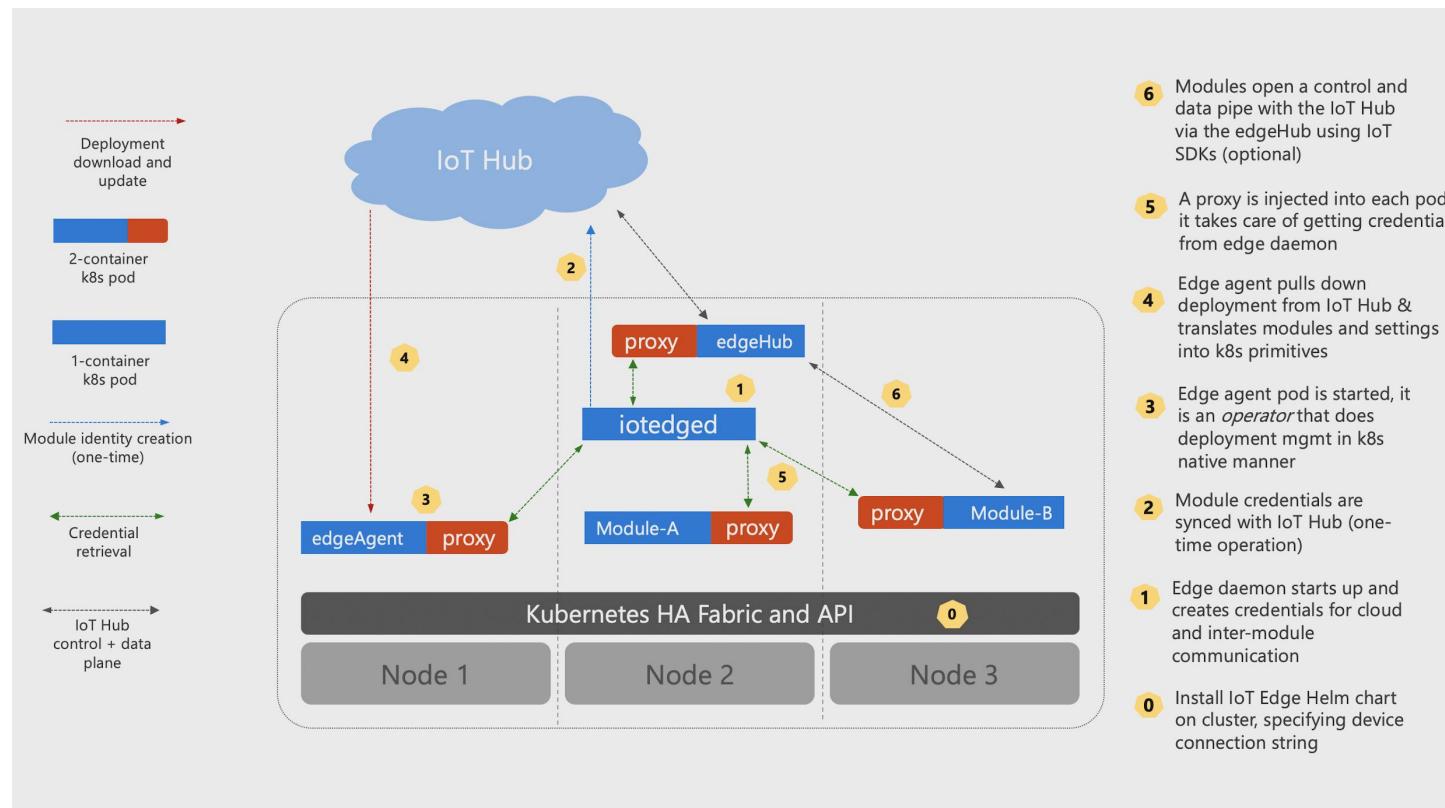
```
apiVersion: containerdays.io/v1
kind: Sensor
metadata:
  name: hamburg
  namespace: germany
spec:
  unit: Celcius
status:
  lastUpdated: 2010-03-26T15:13:42.05Z
  temperature: 28
```

Source: Stefan Schimanski presentation on Extending Kubernetes
<https://www.slideshare.net/sttts/extending-kubernetes-with-customresourcedefinitions>

Another example of a CRD for IoT

Azure IoT Edge

Edge workloads deployed to on premise Kubernetes clusters. Uses [Custom Resource Definitions](#) (CRDs), with a [Controller](#) (IoT Edge Agent) that reconciles cloud managed desired state with the local device state



See details here:
github.com/Azure-Samples/iotedge-gateway-on-kubernetes

Emerging Use case for Kubernetes

What if we view it as a customizable general purpose platform for managing “things”

Originally developed for specific use case of managing containerized apps in a data center

Kubernetes really is a resource management platform

Provides an API that accepts descriptions of desired states to drive controllers/operators



What already invented “wheels” could we re-use?

Popular, well architected API

Tremendous and growing collateral toolchain.. metrics, monitoring, more...



KubeCon



CloudNativeCon

North America 2019

Intermission

5 minutes - birds of a feather



KubeCon



CloudNativeCon

North America 2019

Deep Dive

Futures: Could the service mesh concept be extended beyond TCP, HTTP(s)?

Whether a location has devices, apps or both:
configuration and software update management
is needed

This applies to devices, but also device
connectivity

Photo: Rampion Offshore Wind Farm, United Kingdom
Nicholas Doherty on [Unsplash](#)

Probably not the most efficient
“over the air” update process

Photo: [Wikipedia / Tobias Klenze](#) / CC-BY-SA 4.0.



Edge compute

LTE 5 impact



“By 2022, more than half of enterprises-generated data will be created and processed outside of data centers, and outside of cloud. “ Gartner Dec 2018

<https://www.gartner.com/en/newsroom/press-releases/2018-12-03-gartner-says-the-future-of-it-infrastructure-is-always-on-always-available-everywhere>

Service Mesh =

a data plane +

a control plane used to configure rules that control the data plane

Features:

load balancing, traffic management, routing, health monitoring (maybe not just network health), security policies, service and user authentication, and protection against intrusion and DDoS attacks + offload code for all this from the app developer. Maybe, service discovery add traceability

Sits between application and transport, manage from an application connection perspective, not from the low level infrastructure perspective



github.com/networkservicemesh/networkservicemesh

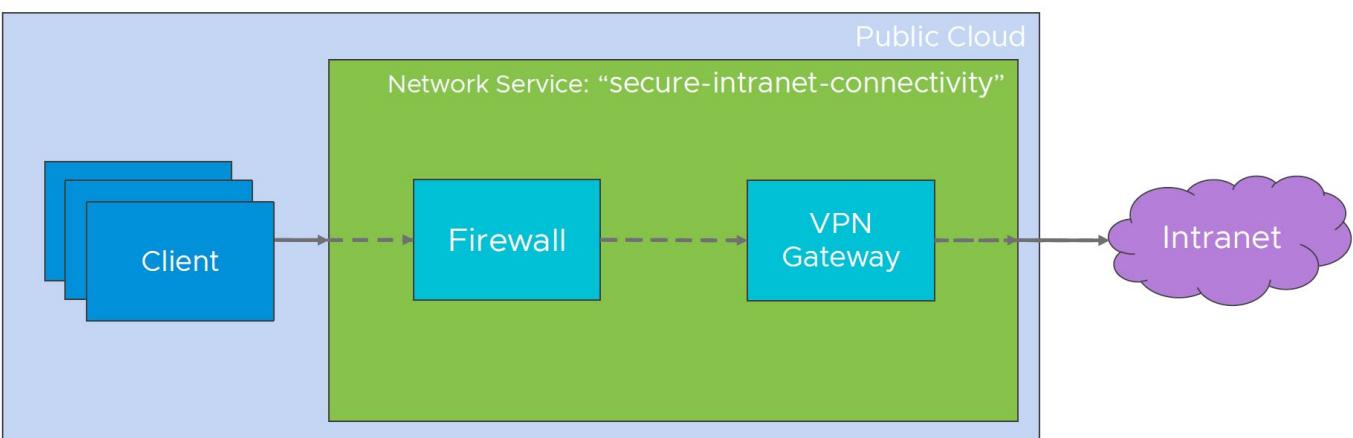
Network Service Mesh

- Service mesh for L2/L3 payloads
- On-demand, dynamic, negotiated connections
- gRPC API to publish and consume Network Services
- Without changes to Kubernetes
- Works with any CNI
- Workload-To-Workload granular level of connectivity
- Loosely coupled heterogeneous network configurations

Network Service Mesh

Define a Network Service

- Specify type of payload
- Source and destination selection
- Service composition



```
apiVersion: networkservicemesh.io/v1
kind: NetworkService
metadata:
  name: secure-intranet-connectivity
spec:
  payload: IP
  matches:
    - match:
        sourceSelector:
          app: firewall
      route:
        - destination:
            destinationSelector:
              app: vpn-gateway
    - match:
        route:
          - destination:
              destinationSelector:
                app: firewall
```

Edge networking

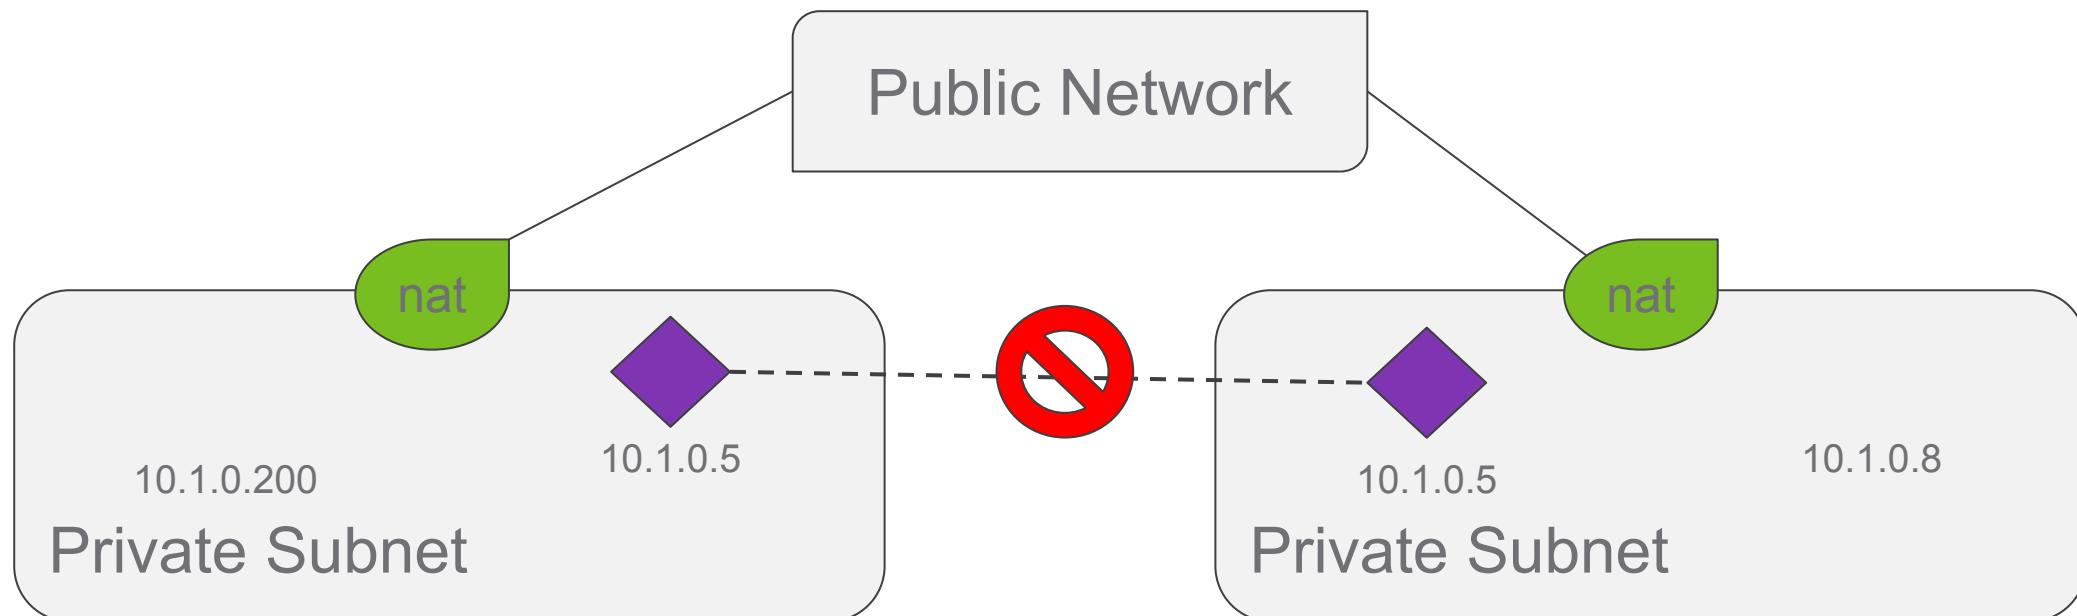


KubeCon

CloudNativeCon

North America 2019

- Hybrid cloud, microservice architecture, agile integration, etc.
 - Not client/server
 - Services/processes want to be deployable and addressable everywhere (north/south/east/west)
- Edge computing - Lots of private subnetworks



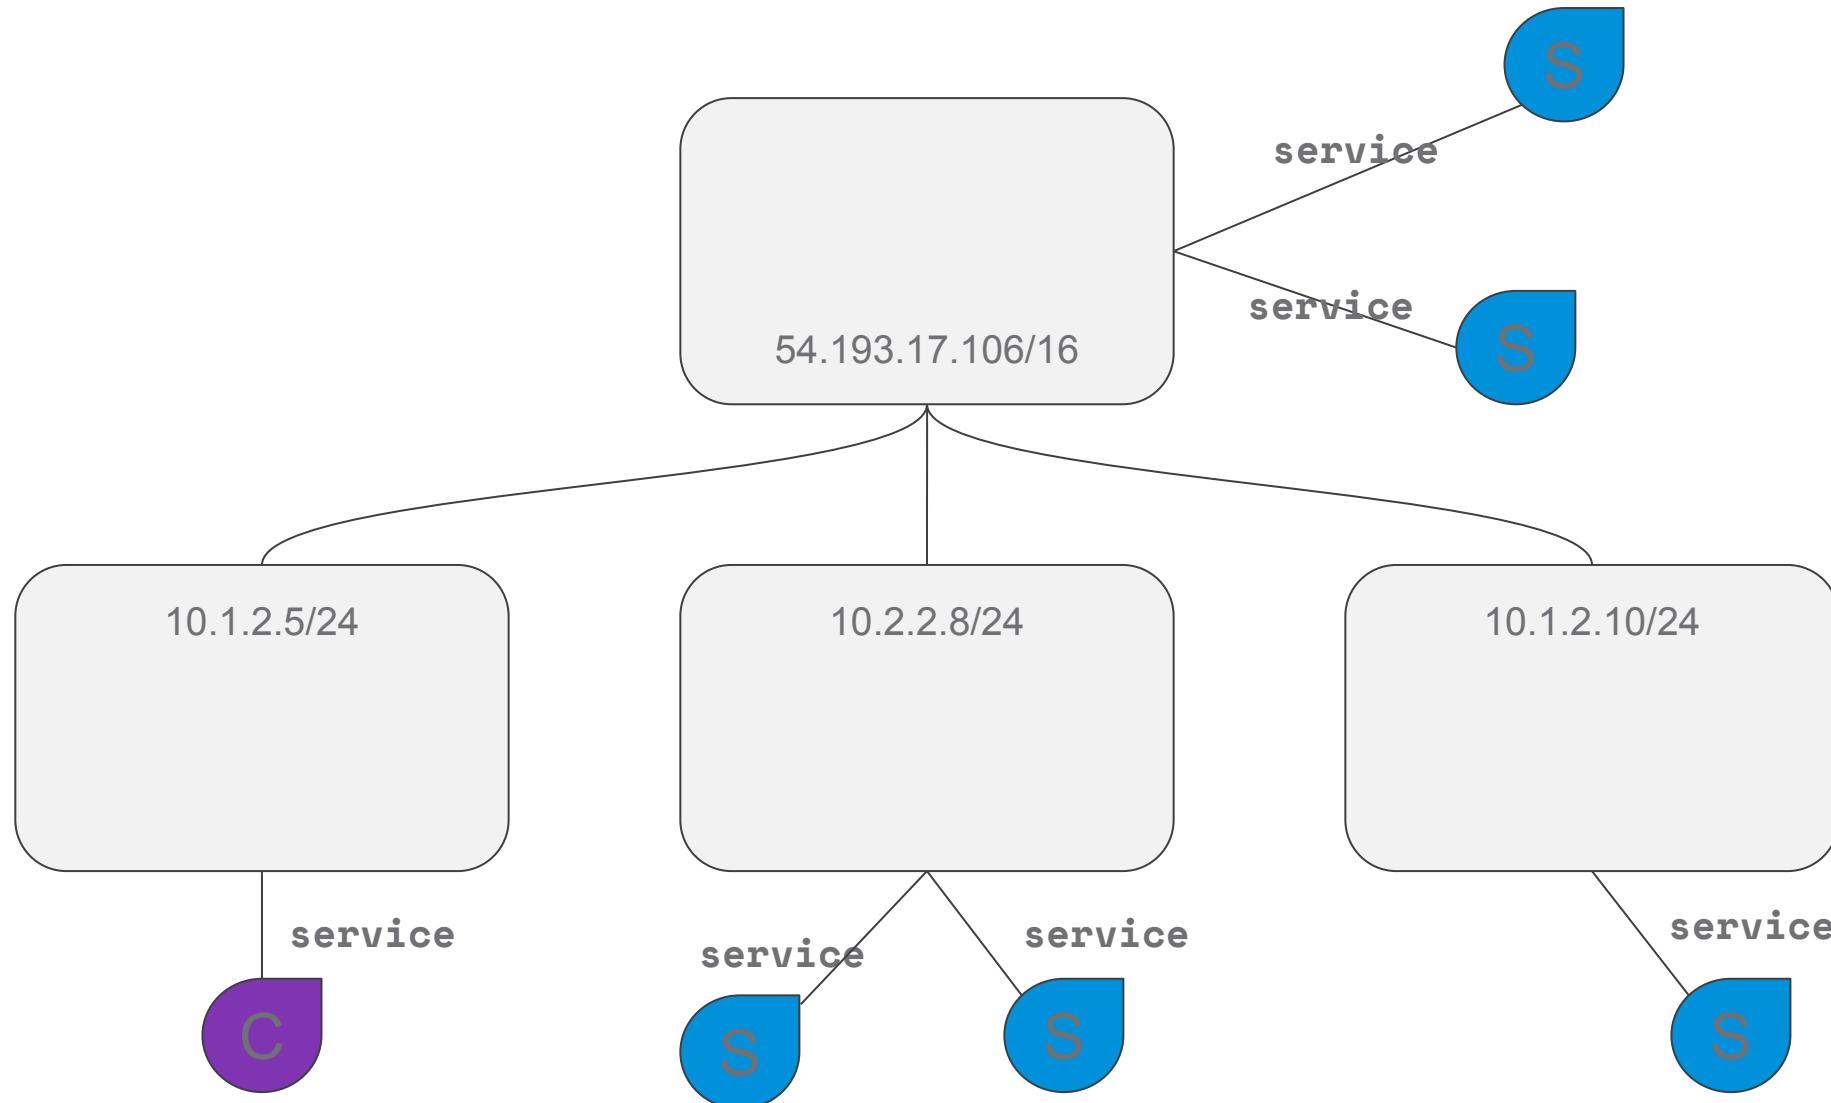
Application Layer Addressing



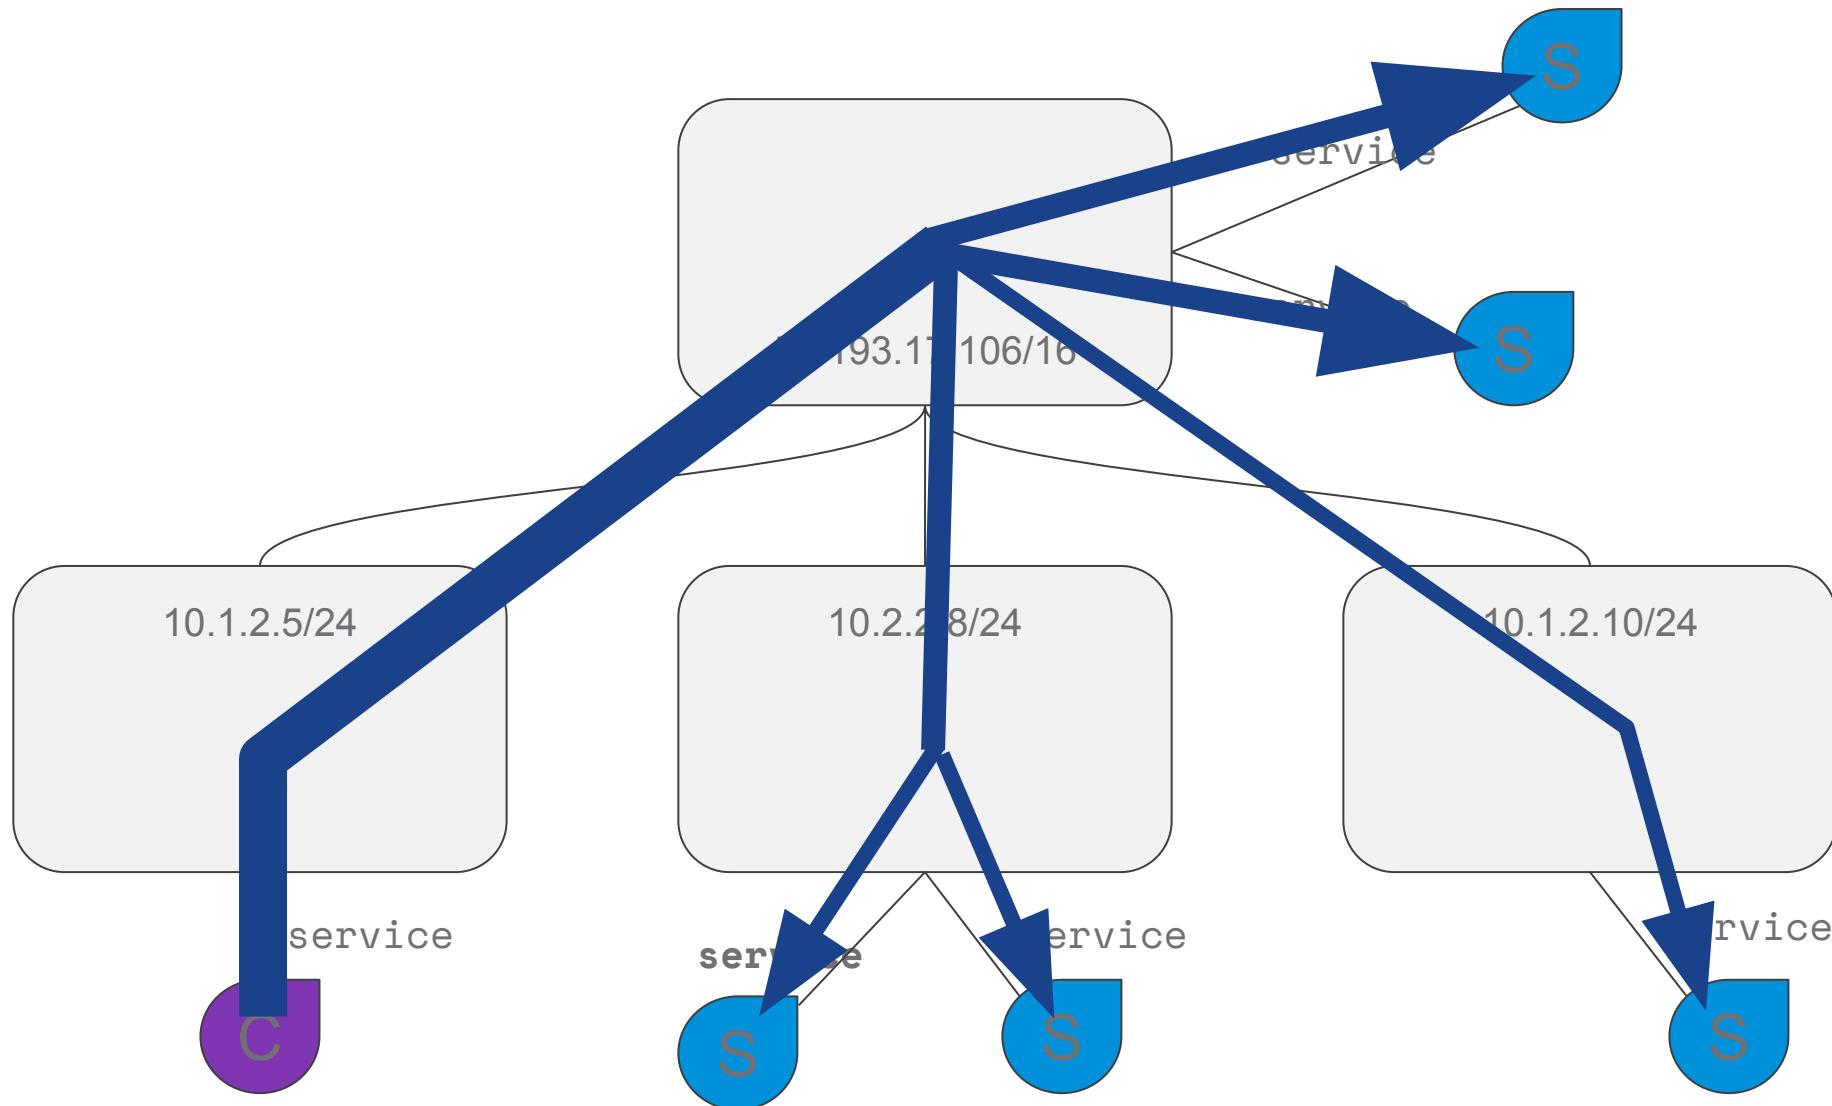
KubeCon

CloudNativeCon

North America 2019



Application Layer Addressing



Secure multi-site communication

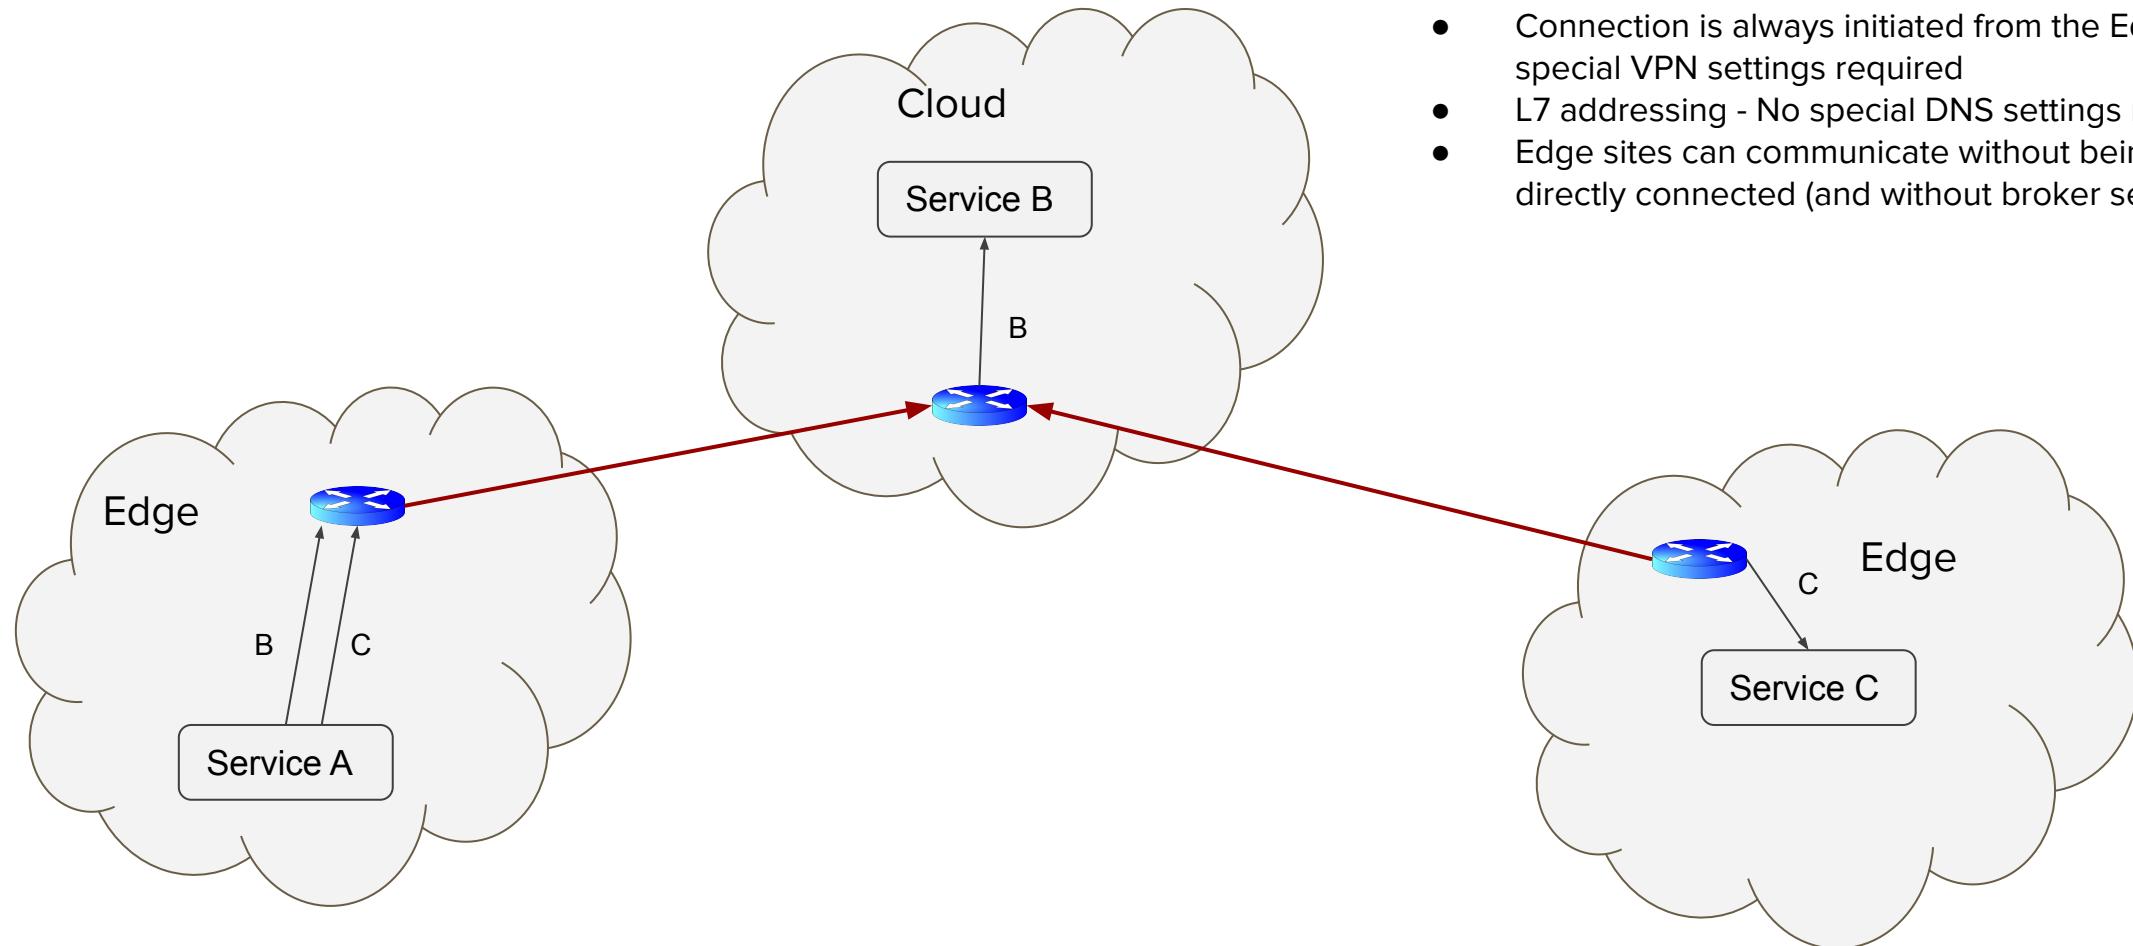


KubeCon



CloudNativeCon

North America 2019



Virtual Application Networks

- Extends a private L7 network across a public network
- Enables users to send and receive data across a shared network as if their computing devices were directly connected to the private network
- The network association is defined by the application's L7 addressing
- E.g. Applications can communicate with other applications through their addressing endpoints independently of the LAN or Subnet the computing device is operating on
- VAN's are built on top of Application Routers



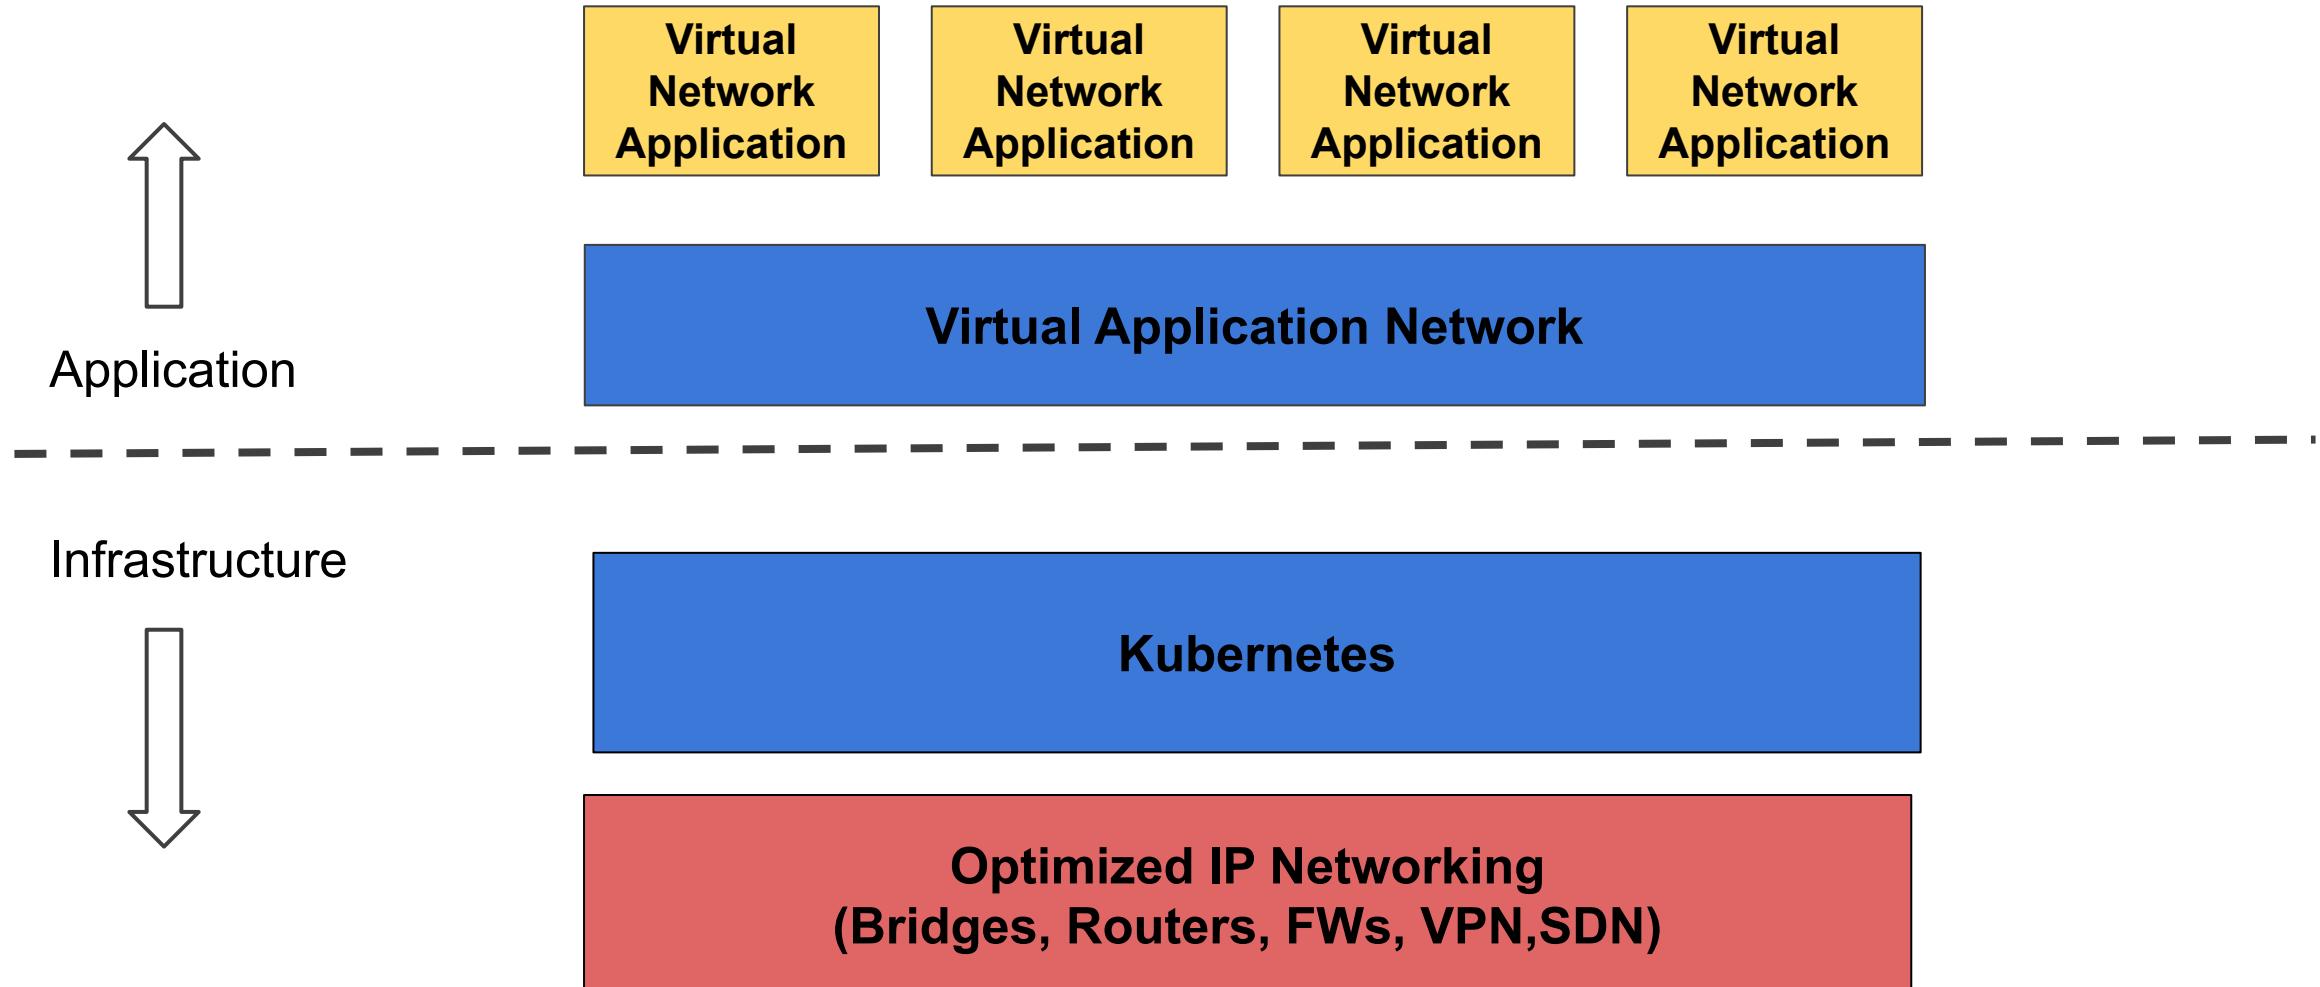
KubeCon



CloudNativeCon

North America 2019

Virtual Application Network “Layer”



Mapping of endpoint Names to VAN Addresses

Mapping of VAN Addresses to Locations

Network of Routers in each Location

TCP/IP

Implications of Application Addressing



- Security
 - Access control for addresses - at the service/process/business resolution
 - Locked-down network membership - Mutual TLS for inter-site connections
 - Cross-cluster applications not exposed via Kube networking
 - Public exposure limited to ingress
 - Trusted and untrusted edges
- Management
 - Metrics collected at business resolution



- Operational Ease

- Easy to deploy in a multi-cluster network
- No advanced networking (SDN, VPNs, Tunnels, Firewall rules, etc.)
- No need for elevated or admin privileges
- No problem with overlapping CIDR subnets or mixes of IPv4 and IPv6
- No single point of failure - use redundant topology

- Not just for messaging

- Proxy maps HTTP, TCP, UDP, etc. to AMQP

- <http://skupper.io>

- Examples, demo-videos, etc.
- New, emerging project



KubeCon



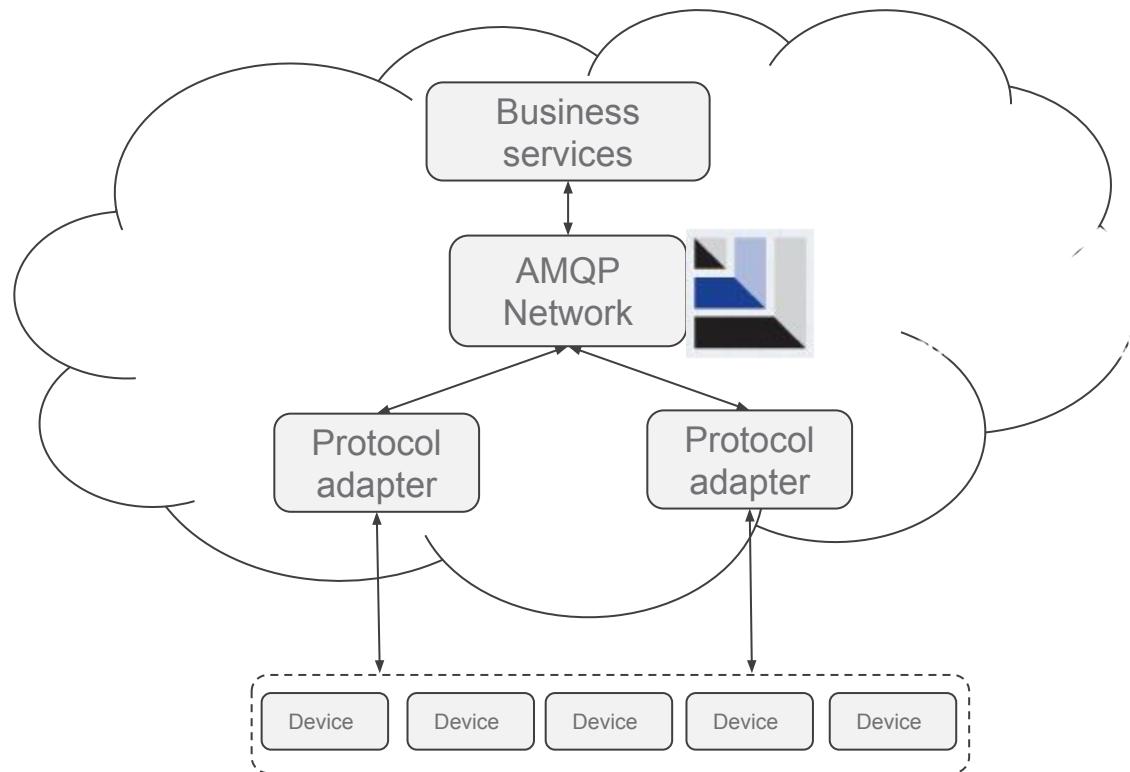
CloudNativeCon

North America 2019

Deep Dive

Futures: Upcoming development, bringing it all together!

Eclipse Hono in the cloud





KubeCon



CloudNativeCon

North America 2019

Is cloud obsolete?

Way forward

- Cloud is not obsolete
- Cloud IoT platforms still needed
 - Business applications
 - Long term data storage
- Work on distributed Edge deployments for IoT services

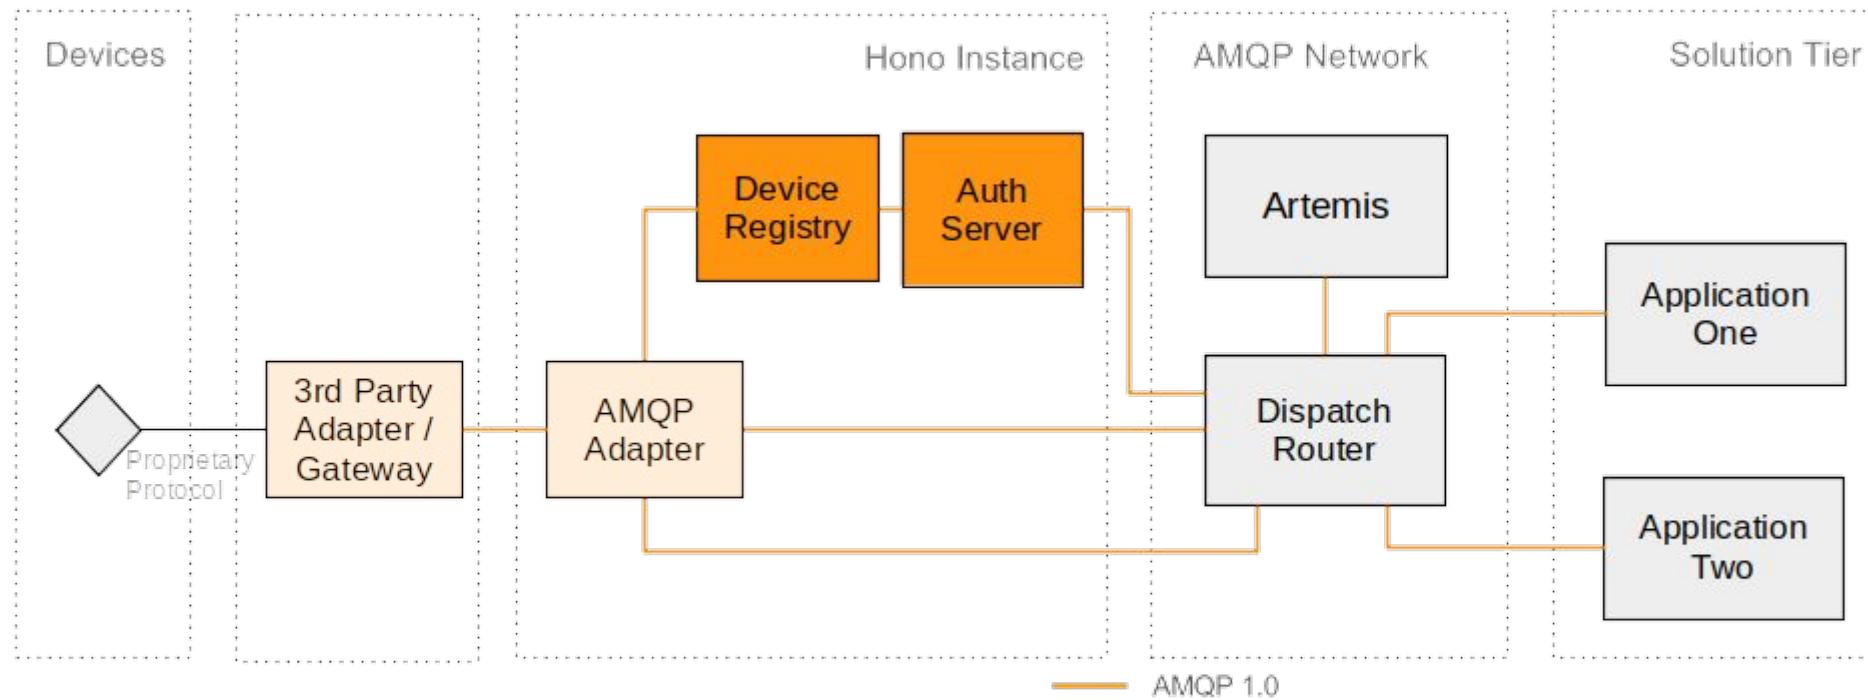


KubeCon

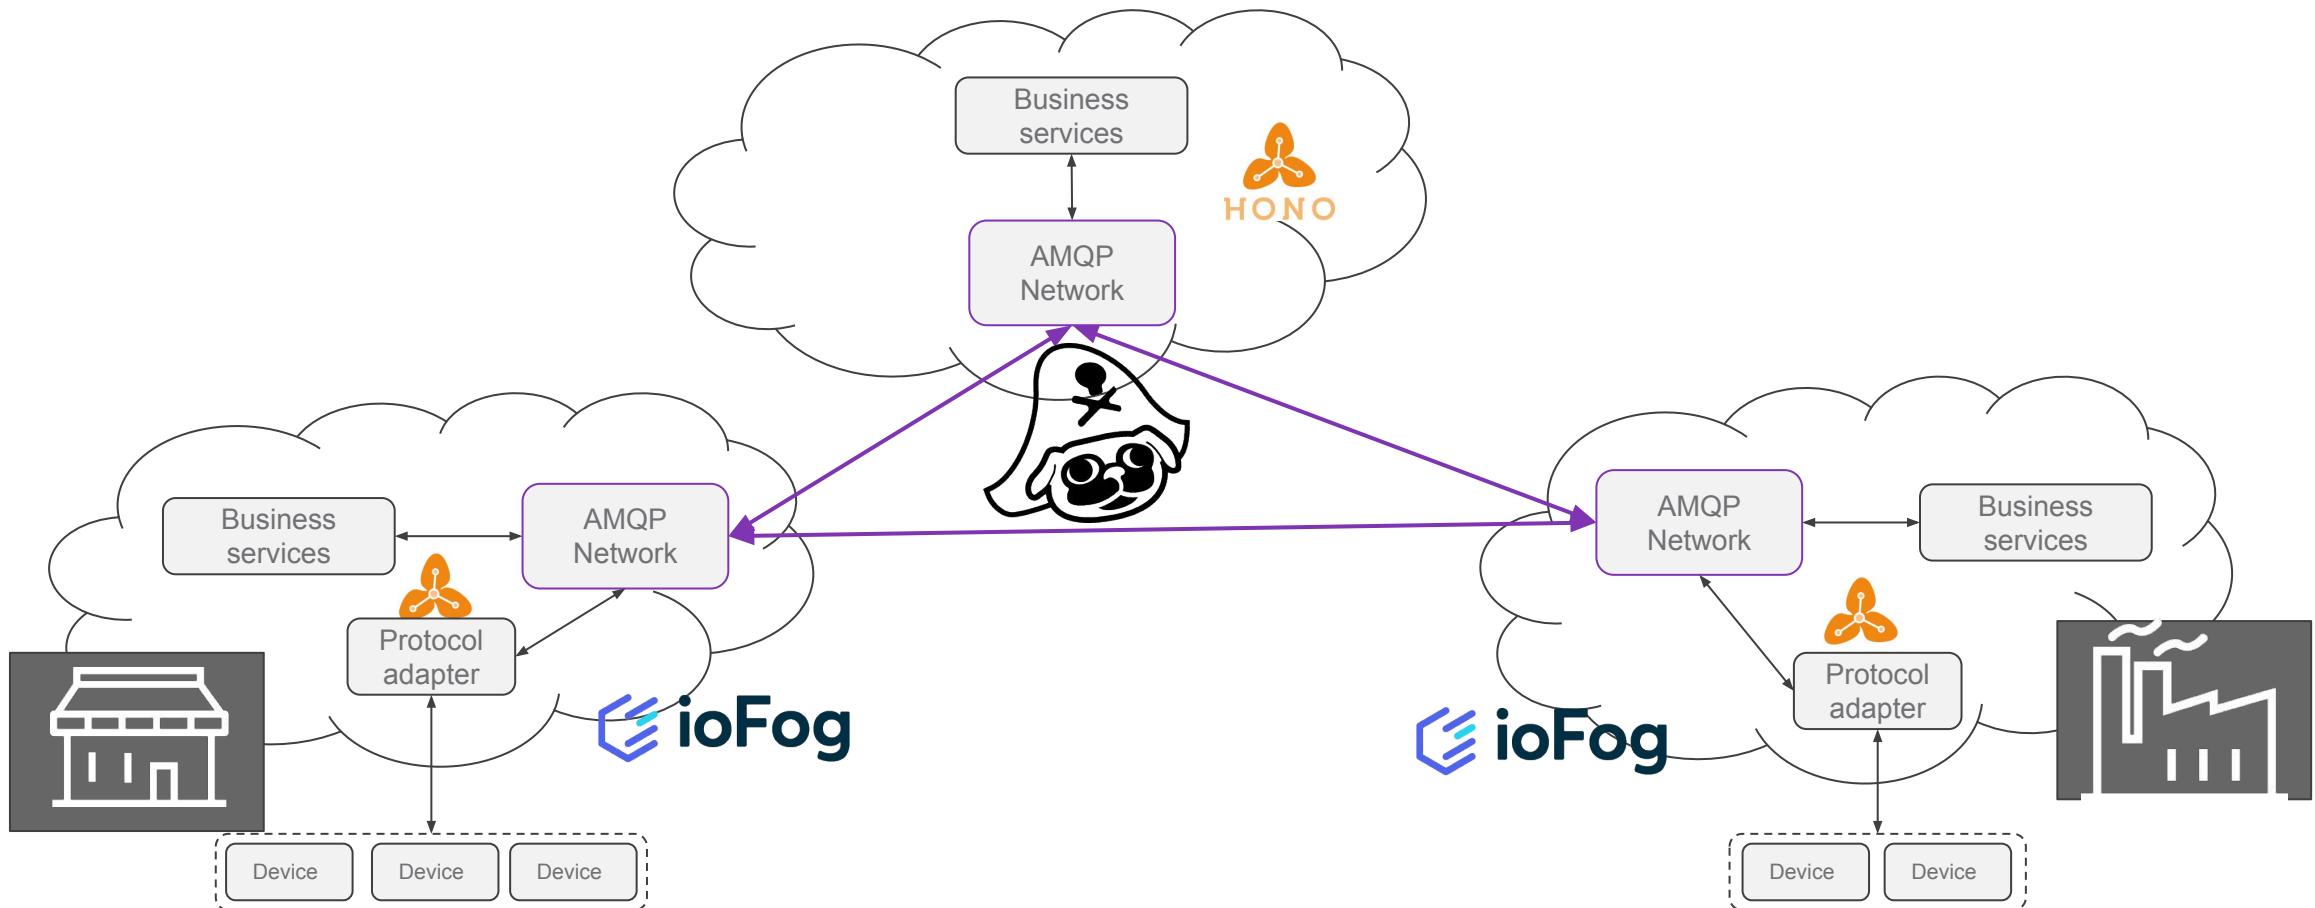


CloudNativeCon

North America 2019



Eclipse Hono on the Edge



Full integration

- Provide development continuum from field to cloud via edge
- Better integration
 - Platforms
 - KubeEdge
 - ioFog
 - Communication
 - Hono
 - Skupper



KubeCon



CloudNativeCon

North America 2019

Demonstrations

- Demonstration: Kubernetes management of an edge application using a specialized protocol
- Demonstration: Use a device gateway with Kubernetes

How to get involved with the IoT Edge Working Group

Learn more.....



Regular Work Group Meeting:

USA WG Meeting Wednesday 9am PT, every 4 weeks, next on December 4

APAC WG meeting Wednesday 5 UTC every 4 weeks, next on December 17

- [Meeting notes and agenda](#)

Link to join the group

- [groups.google.com/forum/#!forum/kubernetes-wg-iot-edge](#)

Link to join Slack

- <https://kubernetes.slack.com/messages/wg-iot-edge>

White Paper

- <http://bit.ly/iot-edge-whitepaper>



KubeCon



CloudNativeCon

North America 2019

Thank You